

# CSE 210 Operating Systems



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Module3 - Process Synchronization and Deadlocks

The Critical-Section Problem- Peterson's Solution- Semaphores, Classic Problems of Synchronization, Monitors.

Introduction to Deadlocks, Necessary conditions for deadlock, deadlock Characterization, Deadlock Prevention- Deadlock Avoidance- Deadlock detection & recovery from Deadlock.

# Process Synchronization

- Background, Race Condition
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Semaphores
- Classic Problems of Synchronization
- Monitors
- Synchronization Examples

# Process Sync. Background

- Processes can execute concurrently
  - May be interrupted at any time, partially completing execution
- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Illustration of the problem:  
Suppose that we wanted to provide a solution to the consumer-producer problem that fills ***all*** the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers.



- Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Producer

```
while (true) {  
    /* produce an item in next_produced */  
  
    while (counter == BUFFER_SIZE); /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Consumer

```
while (true) {  
    while (counter == 0); /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Race Condition

- **counter++** may be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- **counter--** may be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Race Condition

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute <code>register1 = counter</code>	{register1 = 5}
S1: producer execute <code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute <code>register2 = counter</code>	{register2 = 5}
S3: consumer execute <code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute <code>counter = register1</code>	{counter = 6 }
S5: consumer execute <code>counter = register2</code>	{counter = 4}

- The value of **count** may be either 4 or 6, where the correct result should be 5.

# Race Condition

- If both the producer and consumer attempt to update the buffer concurrently, the assembly language statements may get interleaved.
- Interleaving depends upon how the producer and consumer processes are scheduled.

# Race Condition

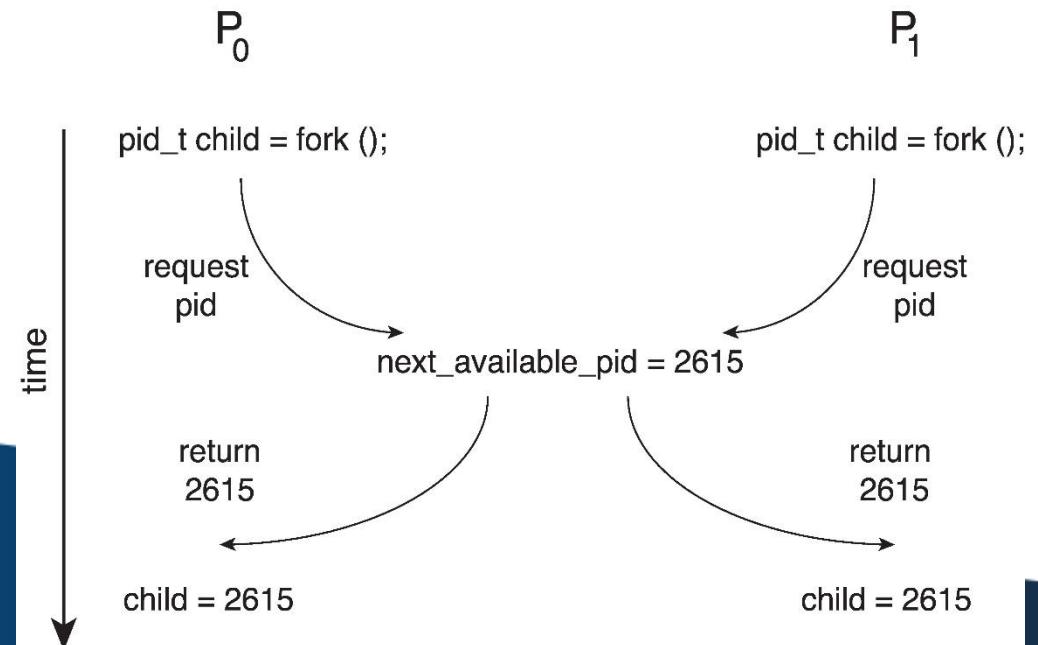
- The statements
  - **counter++;**
  - **counter--;**
  - must be performed *atomically*.
- Atomic operation means an operation that completes in its entirety without interruption.

# Race Condition

- **Race condition:** The situation where several processes access – and manipulate shared data concurrently. The final value of the shared data depends upon which process finishes last.
- To prevent race conditions, concurrent processes must be **synchronized**.

# Race Condition

- Processes  $P_0$  and  $P_1$  are creating child processes using the `fork()` system call
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (pid)

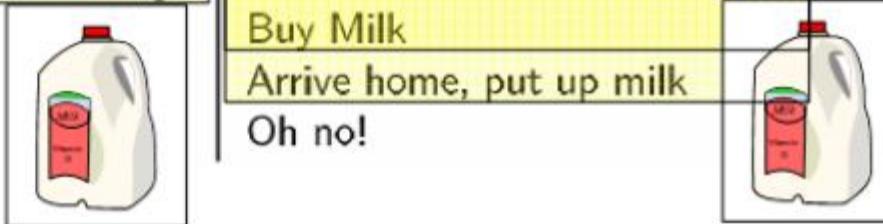


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



<i>time</i>	<i>You</i>	<i>Your Roommate</i>
3:00	Arrive home	
3:05	Look in fridge, no milk	
3:10	Leave for grocery	
3:15		Arrive home
3:20	Arrive at grocery	Look in fridge, no milk
3:25	Buy milk	Leave for grocery
3:35	Arrive home, put milk in fridge	
3:45		Buy Milk
3:50		Arrive home, put up milk
3:50		Oh no!



### Example of Race: The ATM “Withdraw at the same time from a joint account” Problem

Balance = \$300

ATM -1

\$300 ← Check Balance

??? ← Withdraw \$300

ATM-2

Check Balance → \$300

Withdraw \$300 → ???

# Critical Section Problem

- Consider system of  $n$  processes  $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code
  - Process may be changing common variables, updating table, writing file, etc
  - When one process in critical section, no other may be in its critical section
- ***Critical section problem*** is to design protocol to solve this
- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Critical Section

- General structure of process  $P_i$

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

# Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non-preemptive

- **Preemptive** – allows preemption of process when running in kernel mode
- **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU
  - Essentially free of race conditions in kernel mode

# Initial Attempt to Solve the problem

- Only 2 processes,  $P_0$  and  $P_1$
- General structure of process  $P_i$  (other process  $P_j$ )

**do {**

*entry section*

critical section

*exit section*

reminder section

**} while (1);**

- Processes may share some common variables to synchronize their actions.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Algorithm 1

## Shared variables:

```
int turn;  
initially turn = 0
```

**turn - i**  $\Rightarrow$   $P_i$  can enter its critical section

Process  $P_i$

```
do {  
    while (turn != i) ;  
        critical section  
    turn = j;  
        reminder section  
} while (1);
```

Process  $P_j$

```
do {  
    while (turn != j) ;  
        critical section  
    turn = i;  
        reminder section  
} while (1);
```

Satisfies mutual exclusion, but not progress

$P_i, P_j, P_i, P_j \dots$  strict alternation of processes

$P_i$  leaves,  $P_j$  busy with long I/O,  $P_i$  comes back to CS-entry;

No one in the CS, but  $P_i$  has to wait until  $P_j$  to come to the CS

What if  $P_j$  never comes back to CS ????



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Algorithm 2

## Shared variables

**boolean flag[2];**

initially **flag [0] = flag [1] = false.**

**flag [i] = true**  $\Rightarrow P_i$  ready to enter its critical section

Process  $P_i$

**do {**

**flag[i] := true;**

**while (flag[j]);**

critical section

**flag [i] = false;**

remainder section

**} while (1);**

Process  $P_j$

**do {**

**flag[j] := true;**

**while (flag[i]);**

critical section

**flag [j] = false;**

remainder section

**} while (1);**

Satisfies mutual exclusion and progress, but not bounded waiting requirement.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Peterson's Solution

- Not guaranteed to work on modern architectures! (But good algorithmic description of solving the problem)
- Two process solution
- Assume that the `load` and `store` machine-language instructions are atomic; that is, cannot be interrupted

# Two process solution

- Assume that the LOAD and STORE instructions are atomic
- The two processes share two variables:
  - int turn;
  - Boolean flag[2]      Initially: flag[0]=flag[1]= false
- The variable turn indicates whose turn it is to enter the critical section.
- The flag array is used to indicate if a process is ready to enter the critical section.  
 $\text{flag}[i] = \text{true}$  implies that process  $P_i$  is ready!

# Algorithm for Process $P_i, P_j$

## The Critical Section Problem:

		flag	false	false	turn	i	
		i	j				
process $P_i$		process $P_j$					
	do {						
entry section		flag[i] = true; turn = j; while (flag[j] && turn == j);					
		<i>critical section</i>					
exit section		flag[i] = false;					
		<i>remainder section</i>					
	} while (1);						
		do {					
		flag[j] = true; turn = i; while (flag[i] && turn == i);					
		<i>critical section</i>					
		flag[j] = false;					
		<i>remainder section</i>					
		}` while (1);					

# Algorithm 3 - cont.

- Claim: Mutual exclusion is preserved
- Proof:
  - Assume Pi and Pj in CS at the same time
  - flag[i] & flag[j] == True
  - (I) flag[j] & turn == j → False → turn = i
  - (II) flag[i] & turn == i → False → turn = j
  - turn cannot be i and j at the same time

# Algorithm 3- cont.

- Claim: Progress is met.
- Proof:
  - Pi stuck at “flag[j] == true & turn == j” while loop
  - (case 1) Pj is not ready to enter
    - flag[j] == false → Pi can enter
  - (case 2) Pj is ready to enter
    - Pj : set flag[j] to true and at its while
    - Observation: turn == i XOR turn == j
    - (case 2.1) turn == i → Pi will enter
    - (case 2.2) turn == j → Pj will enter
      - (case 2.2.1) Pj leaves CS; sets flag[j] to false → Pi enters
      - (case 2.2.2) Pj leaves CS; sets flag[j] to true ;
        - then sets turn to i → Pi enters
        - (Pi at while cannot change turn)
- Pi (Pj) will enter the CS (progress) after at most one entry by Pj (Pi) (bounded waiting)

# Peterson's Algorithm

- Less complicated than Dekker's Algorithm
  - Still uses busy waiting.
  - Requires fewer steps to perform mutual exclusion primitives
  - Easier to demonstrate its correctness



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Synchronization Hardware

- Many systems provide hardware support for implementing the critical section code.
- All solutions below based on idea of **locking**
  - Protecting critical regions via locks
- Uniprocessors – could disable interrupts
  - Currently running code would execute without preemption
  - Generally too inefficient on multiprocessor systems
    - Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
  - **Atomic** = non-interruptible
    - Either test memory word and set value
    - Or swap contents of two memory words



# Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# test\_and\_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to “TRUE”.

# Solution using test\_and\_set()

- Shared Boolean variable lock, initialized to FALSE
- Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
        /* critical section */  
    lock = false;  
        /* remainder section */  
} while (true);
```

# compare\_and\_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
  
    return temp;  
}
```

1. Executed atomically
2. Returns the original value of passed parameter “value”
3. Set the variable “value” the value of the passed parameter “new\_value” but only if “value” ==“expected”. That is, the swap takes place only under this condition.

# Solution using compare\_and\_swap

- Shared integer “lock” initialized to 0;
- Solution:

```
do {
    while (compare_and_swap(&lock, 0, 1) != 0)
        ; /* do nothing */
    /* critical section */
    lock = 0;
    /* remainder section */
} while (true);
```

# Bounded-waiting Mutual Exclusion with test\_and\_set

Boolean waiting[n]; and Boolean lock are set to false.

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)

        key = test_and_set(&lock);

    waiting[i] = false;

    /* critical section */

    j = (i + 1) % n;

    while ((j != i) && !waiting[j])

        j = (j + 1) % n;

    if (j == i)

        lock = false;

    else

        waiting[j] = false;

    /* remainder section */

} while (true);
```



# Mutex Locks

- Previous solutions are complicated and generally inaccessible to application programmers
- OS designers build software tools to solve critical section problem
- Simplest is mutex lock
- Protect a critical section by first **acquire()** a lock then **release()** the lock
  - Boolean variable indicating if lock is available or not
- Calls to **acquire()** and **release()** must be atomic
  - Usually implemented via hardware atomic instructions
- But this solution requires **busy waiting**
  - This lock therefore called a **spinlock**



# acquire() and release()

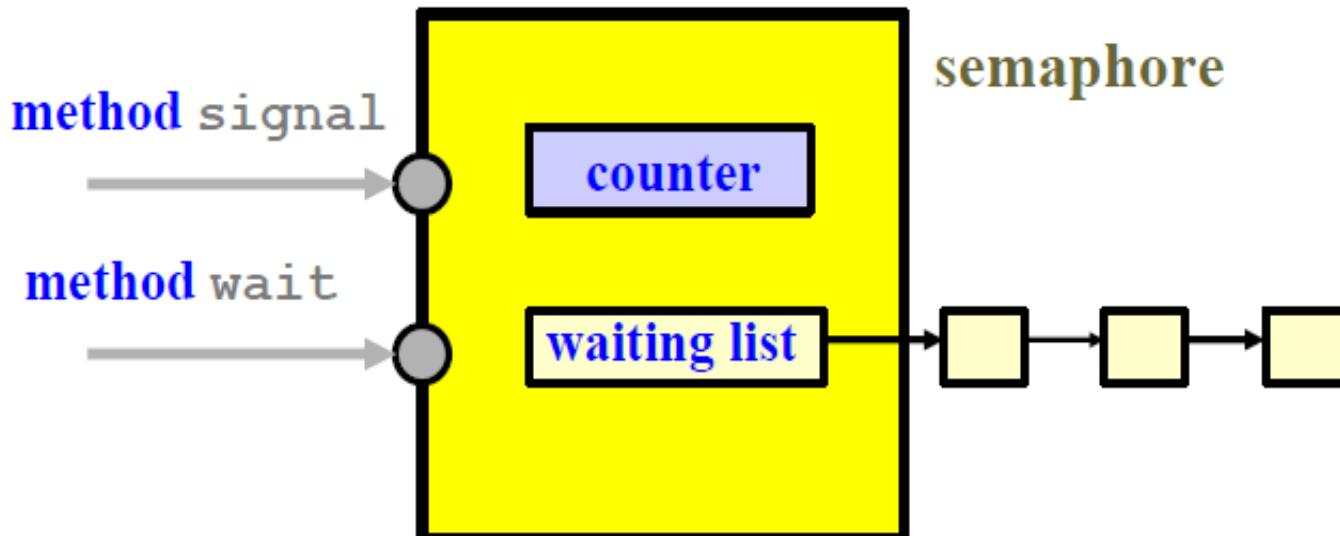
- ```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}  
• release() {  
    available = true;  
}  
• do {  
    acquire lock  
  
    critical section  
  
    release lock  
    remainder section  
} while (true);
```



# Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.

❑ A *semaphore* is an object that consists of a counter, a waiting list of processes and two methods (e.g., functions): signal and wait.



# Semaphore Method: wait

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting list;
        block();
    }
}
```

- After decreasing the counter by 1, if the counter value becomes negative, then
  - ❖ add the caller to the waiting list, and then
  - ❖ block itself.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Semaphore Method: signal

```
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a process P from the waiting list;
        resume(P);
    }
}
```

- After increasing the counter by 1, if the new counter value is not positive, then
  - ❖ remove a process P from the waiting list,
  - ❖ resume the execution of process P, and return

# Important Note: 1/4

```
S.count--;  
if (S.count<0) {  
    add to list;  
    block();  
}  
  
S.count++;  
if (S.count<=0) {  
    remove P;  
    resume (P);  
}
```

- If `S.count < 0`, `abs(S.count)` is the number of waiting processes.
- This is because processes are added to (*resp.*, removed from) the waiting list only if the counter value is  $< 0$  (*resp.*,  $\leq 0$ ).



# Important Note: 2/4

```
S.count--;  
if (S.count<0) {  
    add to list;  
    block();  
}
```

```
S.count++;  
if (S.count<=0) {  
    remove P;  
    resume (P);  
}
```

- The waiting list can be implemented with a queue if FIFO order is desired.
- However, the correctness of a program should not depend on a particular implementation of the waiting list.
- Your program should not make any assumption about the ordering of the waiting list.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Important Note: 3/4

```
S.count--;  
if (S.count<0) {  
    add to list;  
    block();  
}  
  
S.count++;  
if (S.count<=0) {  
    remove P;  
    resume(P);  
}
```

- The caller may be blocked in the call to `wait()`.
- The caller never blocks in the call to `signal()`.  
If `S.count > 0`, `signal()` returns and the caller continues. Otherwise, a waiting process is released and the caller continues. In this case, *two processes continue*.



# The Most Important Note: 4/4

```
S.count--;  
if (S.count<0) {  
    add to list;  
    block();  
}  
  
S.count++;  
if (S.count<=0) {  
    remove P;  
    resume(P);  
}
```

- ❑ **wait()** and **signal()** must be executed ***atomically*** (*i.e.*, as one **uninterruptible** unit).
- ❑ Otherwise, ***race conditions*** may occur.
- ❑ **Homework:** use execution sequences to show race conditions if **wait()** and/or **signal()** is not executed atomically.



# Three Typical Uses of Semaphores

- There are three typical uses of semaphores:
  - ❖ **mutual exclusion:**  
    **Mutex (i.e., Mutual Exclusion) locks**
  - ❖ **count-down lock:**  
    Keep in mind that semaphores have a counter.
  - ❖ **notification:**  
    Indicate an event has occurred.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Use 1: Mutual Exclusion (Lock)

```
semaphore S = 1;           initialization is important
int      count = 0;

Process 1
while (1) {
    // do something entry
    S.wait();
    critical sections
    count++;
    S.signal();
    // do something exit
}

Process 2
while (1) {
    // do something
    S.wait();
    critical sections
    count--;
    S.signal();
    // do something
}
```

- What if the initial value of S is zero?
- S is a *binary semaphore* (similar to a *lock*).



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



## Use 2: Count-Down Counter

semaphore S = 3

Process 1

```
while (1) {  
    // do something  
    S.wait();  
    S.signal();  
    // do something  
}
```

Process 2

```
while (1) {  
    // do something  
    S.wait();  
    S.signal();  
    // do something  
}
```

at most 3 processes can be here!!!

- ❑ After three processes pass through `wait()`, this section is locked until a process calls `signal()`.



PRESIDENCY  
UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013



## Use 3: Notification

```
semaphore S1 = 1, S2 = 0;

    process 1
while (1) {
    // do something
    S1.wait(); notify
    cout << "1";
    S2.signal(); notify
    // do something
}

    process 2
while (1) {
    // do something
    S2.wait(); notify
    cout << "2";
    S1.signal(); notify
    // do something
}
```

- Process 1 uses `S2.signal()` to notify process 2, indicating “I am done. Please go ahead.”
- The output is 1 2 1 2 1 2 .....
- What if both `S1` and `S2` are both 0's or both 1's?



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# CSE 210 Operating Systems



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Module3 - Process Synchronization and Deadlocks

The Critical-Section Problem- Peterson's Solution- Semaphores, Classic Problems of Synchronization, Monitors.

Introduction to Deadlocks, Necessary conditions for deadlock, deadlock Characterization, Deadlock Prevention- Deadlock Avoidance- Deadlock detection & recovery from Deadlock.

# Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
  - Bounded-Buffer Problem
  - Readers and Writers Problem

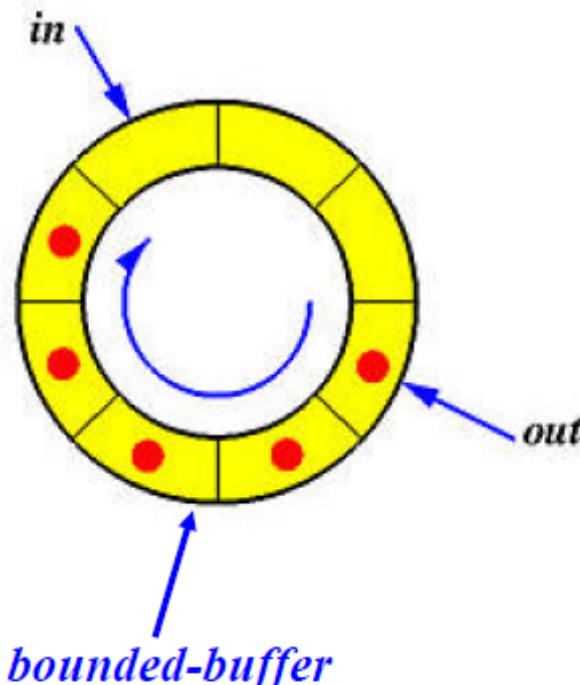


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# The Producer/Consumer Problem



- ❑ Suppose we have a **circular buffer** of  $n$  slots.
- ❑ Pointers **in** (*resp.*, **out**) points to the first **empty** (*resp.*, **filled**) slot.
- ❑ **Producer** processes keep adding info into the buffer
- ❑ **Consumer** processes keep retrieving info from the buffer.

..

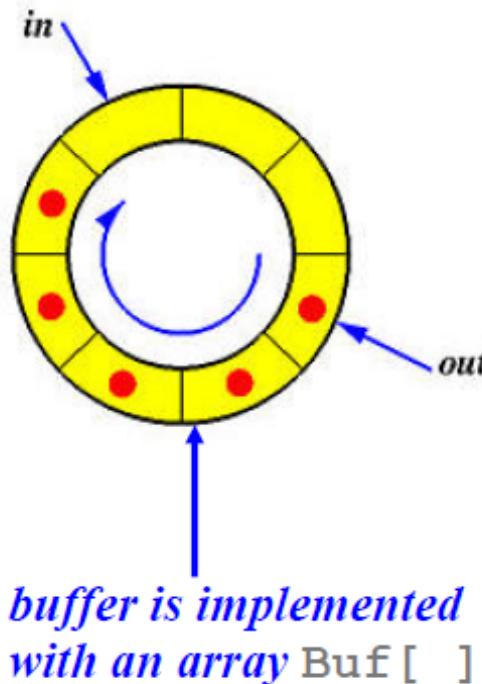


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013

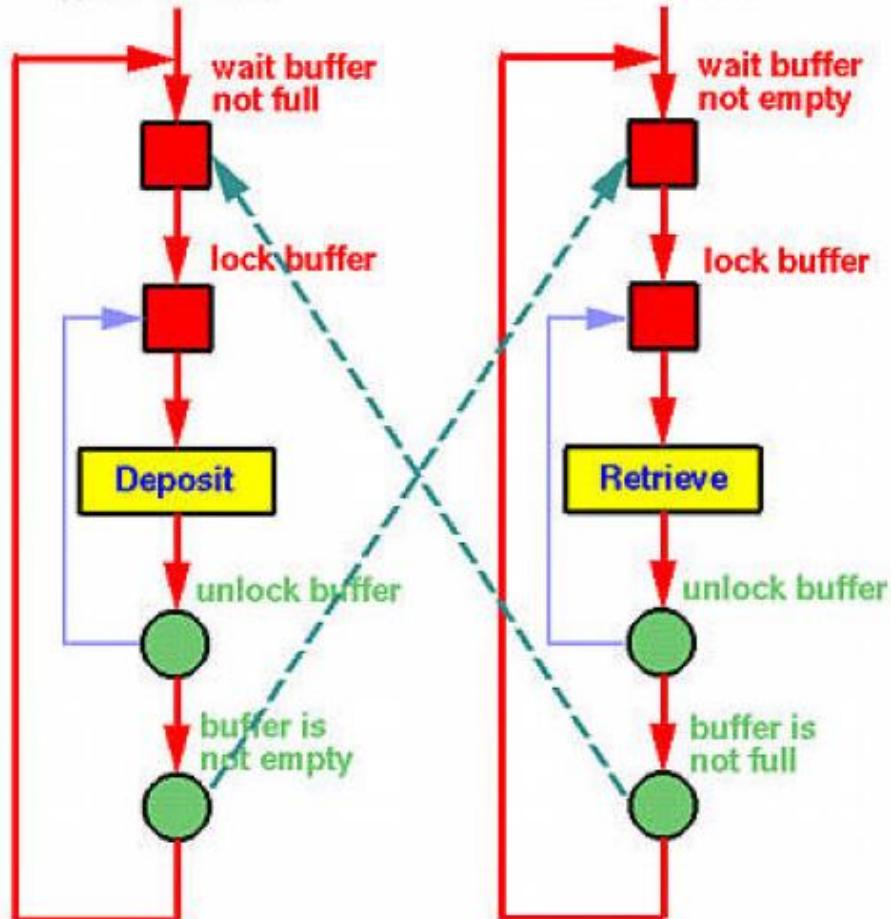


# Problem Analysis



- A producer deposits info into `Buf[in]` and a consumer retrieves info from `Buf[out]`.
- `in` and `out` must be advanced.
- `in` is shared among producers.
- `out` is shared among consumers.
- If `Buf` is full, producers should be blocked.
- If `Buf` is empty, consumers should be blocked.

## *producer*



## *consumer*

- We need a sem. to protect the buffer.
- A second sem. to block producers if the buffer is full.
- A third sem. to block consumers if the buffer is empty.

22



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Solution

no. of slots

```
semaphore NotFull=n, NotEmpty=0, Mutex=1;
```

producer

```
while (1) {  
    NotFull.wait();  
    Mutex.wait();  
    Buf[in] = x;  
    in = (in+1)%n;  
    Mutex.signal();  
    NotEmpty.signal();  
}
```

consumer

```
while (1) {  
    NotEmpty.wait();  
    Mutex.wait();  
    x = Buf[out];  
    out = (out+1)%n;  
    Mutex.signal();  
    NotFull.signal();  
}
```

notifications

critical section



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Question

- ❑ What if the producer code is modified as follows?
- ❑ Answer: a deadlock may occur. Why?

```
while (1) {  
    Mutex.wait();  
    NotFull.wait();  
    Buf[in] = x;  
    in = (in+1)%n;  
    NotEmpty.signal();  
    Mutex.signal();  
}
```

order changed



# The Readers/Writers Problem

- Two groups of processes, **readers** and **writers**, are accessing a shared resource by the following rules:
  - ❖ Readers can **read simultaneously**.
  - ❖ **Only one writer can write at any time.**
  - ❖ When a writer is writing, no reader can read.
  - ❖ If there is any reader reading, all **incoming writers must wait**. Thus, readers have higher priority.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Problem Analysis

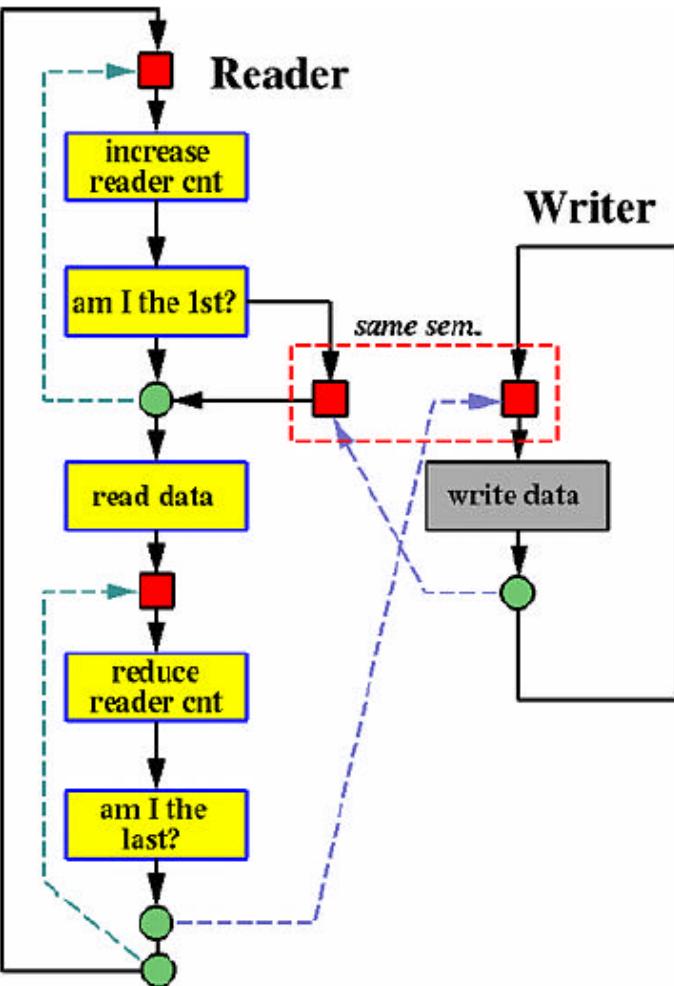
- We need a semaphore to **block** readers if a writer is writing.
- When a writer arrives, it must be able to **know if there are readers reading**. So, a reader count is required which must be protected by a lock.
- This **reader-priority** version has a problem: bounded waiting condition may be violated if readers keep coming, causing the waiting writers no chance to write.



**PRESIDENCY  
UNIVERSITY**

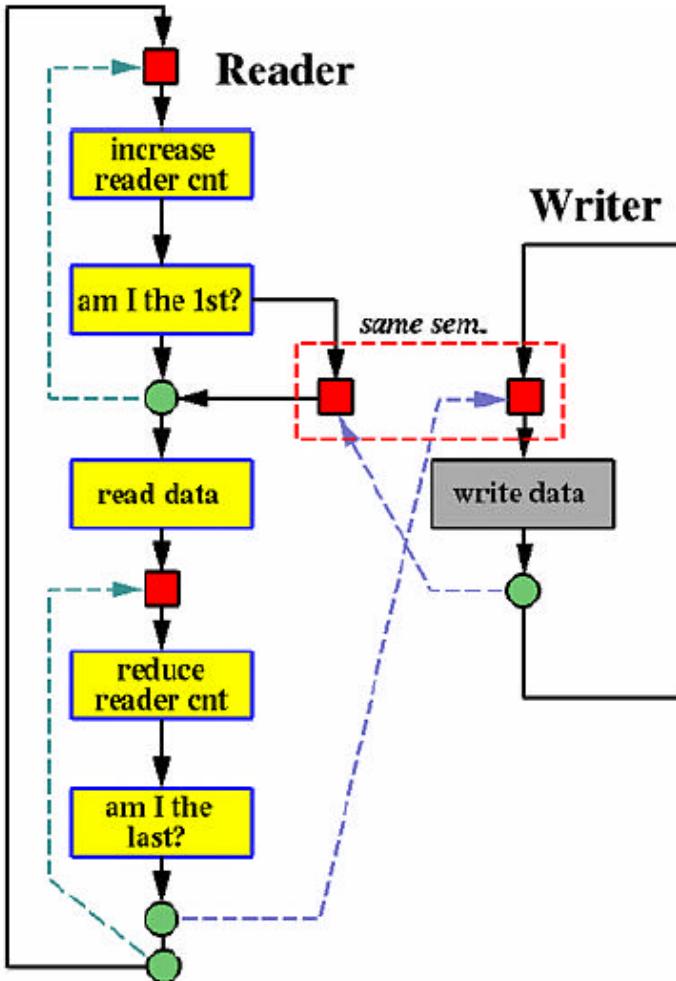
Private University Estd. in Karnataka State by Act No. 41 of 2013





- When a reader comes in, it increases the count.
- If it is the 1<sup>st</sup> reader, waits until no writer is writing,
- Reads data.
- Decreases the counter.
- Notifies the writer that no reader is reading if it is the last.

27



- ❑ When a writer comes in, it waits until no reader is reading and no writer is writing.
- ❑ Then, it writes data.
- ❑ Finally, notifies readers and writers that no writer is in.

28

# Solution

```
semaphore Mutex = 1, WrtMutex = 1;  
int RdrCount;
```

## reader

```
while (1) {  
    Mutex.wait();  
    RdrCount++;  
    if (RdrCount == 1)  
        WrtMutex.wait();  
    Mutex.signal();  
    // read data  
    Mutex.wait();  
    RdrCount--;  
    if (RdrCount == 0)  
        WrtMutex.signal();  
    Mutex.signal();  
}
```

## writer

```
while (1) {  
    WrtMutex.wait();  
    // write data  
    WrtMutex.signal();  
}
```

*blocks both readers and writers*

29



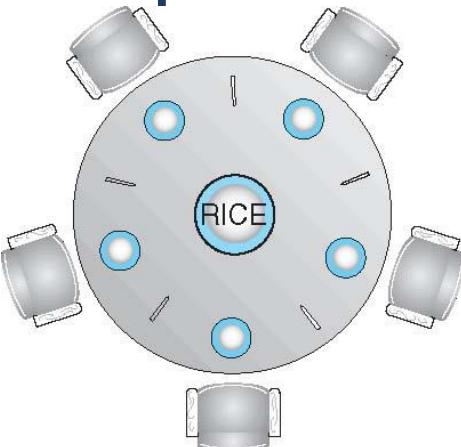
**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Classical Problems of Synchronization

## Dining-Philosophers Problem



- Philosophers spend their lives alternating thinking and eating
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
  - Need both to eat, then release both when done
- In the case of 5 philosophers
  - Shared data
    - Bowl of rice (data set)
    - Semaphore **chopstick [5]** initialized to 1



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Dining-Philosophers Problem Algorithm

- The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i] );  
    wait (chopStick[ (i + 1) % 5] );  
  
        // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
        // think  
  
} while (TRUE);
```

- What is the problem with this algorithm?



- Two neighbours cant eat at same time
- If all philosopher's are hungry then they take Left chopstick lead to deadlock
- If 2 philosophers are faster in taking chopstick ?

### Solution - 1

Take left fork first and then take second fork.

```
do {  
    wait(chopstick[i]);  
    wait(chopstick[(i+1)%5]);  
  
    // Critical Section  
    eat  
  
    signal(chopstick[i]);  
    signal(chopstick[(i+1)%5]);  
  
    think  
}while(true);
```

## Solution - 2

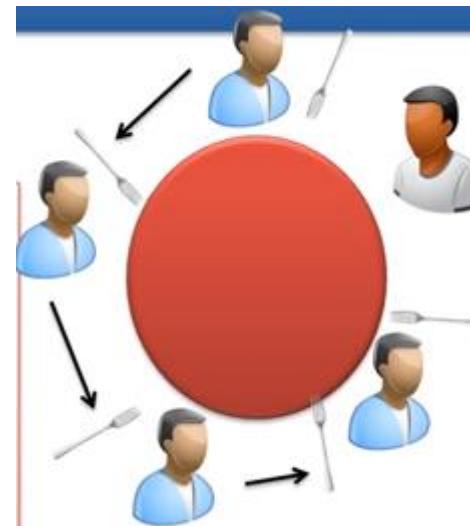
Lefty

```
do {  
    wait(chopstick[i]);  
    wait(chopstick[(i+1)%5];  
  
    // Critical Section  
    eat  
  
    signal(chopstick[i]);  
    signal(chopstick[(i+1)%5];  
  
    think  
}while(true);
```

Righty

```
do {  
    wait(chopstick[(i+1)%5];  
    wait(chopstick[i]);  
  
    // Critical Section  
    eat  
  
    signal(chopstick[(i+1)%5];  
    signal(chopstick[i]);  
  
    think  
}while(true);
```

- Problem - Starvation



**PRESIDENCY  
UNIVERSITY**

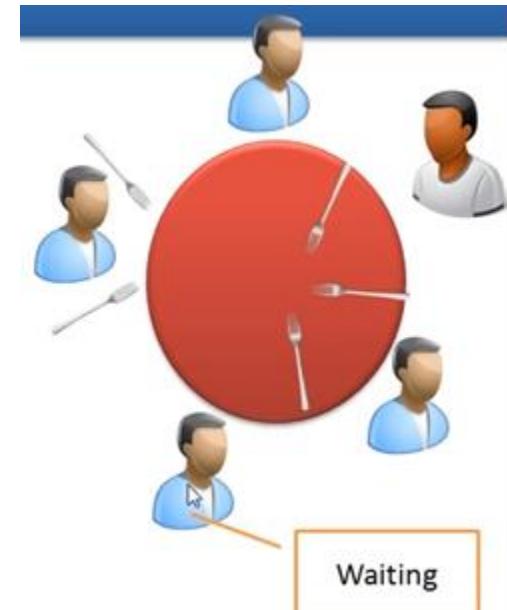
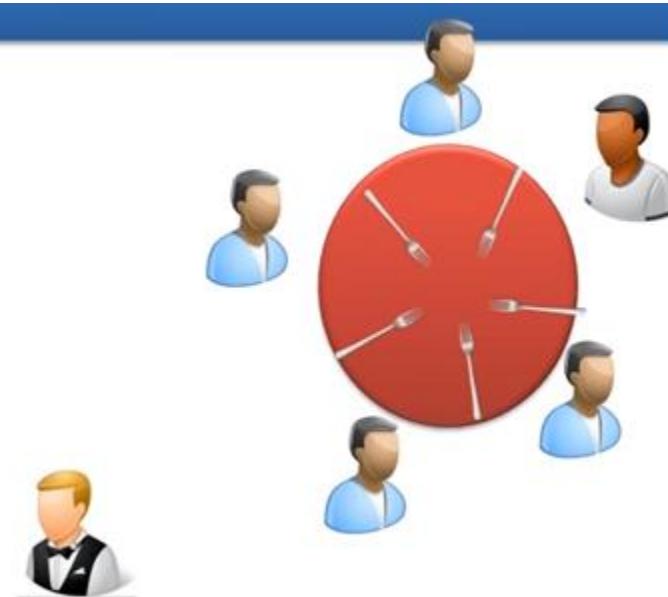
Private University Estd. in Karnataka State by Act No. 41 of 2013



### Solution - 3

#### Use of Arbitrator

```
do {  
    wait(mutex)  
  
    wait(chopstick[i]);  
    wait(chopstick[(i+1)%5];  
    signal(mutex)  
    // Critical Section  
    eat  
  
    signal(chopstick[i]);  
    signal(chopstick[(i+1)%5];  
    think  
}while(true);
```



- This is proposed by Dijkstra



# Tannenbaum solution

## Solution with Semaphores

Uses **N semaphores** ( $s[1], s[2], \dots, s[N]$ ) all initialized to 0, and a mutex  
Philosopher has 3 states: HUNGRY, EATING, THINKING

*A philosopher can only move to EATING state if neither neighbor is eating*

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    unlock(mutex);  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY && state[LEFT] != EATING && state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

```

void philosopher(int i){
    while(TRUE){
        think();
        → take_forks(i);
        eat();
        put_forks();
    }
}

```

```

void take_forks(int i){
    lock(mutex);
    state[i] = HUNGRY;
    test(i);
    unlock(mutex);
    down(s[i]);
}

```

```

void put_forks(int i){
    lock(mutex);
    state[i] = THINKING;
    test(LEFT);
    test(RIGHT)
    unlock(mutex);
}

```

```

void test(int i){
    if (state[i] = HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING){
        state[i] = EATING;
        up(s[i]);
    }
}

```

|           | P1 | P2 | P3 | P4 | P5 |
|-----------|----|----|----|----|----|
| state     | T  | T  | T  | T  | T  |
| semaphore | 0  | 0  | 0  | 0  | 0  |



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    unlock(mutex);  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

|           | P1 | P2 | P3 | P4 | P5 |                                                                                      |
|-----------|----|----|----|----|----|--------------------------------------------------------------------------------------|
| state     | T  | T  | H  | T  | T  |  |
| semaphore | 0  | 0  | 0  | 0  | 0  |                                                                                      |

# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT)  
    unlock(mutex);  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

|           | P1 | P2 | P3 | P4 | P5 |
|-----------|----|----|----|----|----|
| state     | T  | T  | E  | T  | T  |
| semaphore | 0  | 0  | 1  | 0  | 0  |



```

void philosopher(int i){
    while(TRUE){
        think();
        take_forks(i);
        eat();
        put_forks();
    }
}

```

```

void take_forks(int i){
    lock(mutex);
    state[i] = HUNGRY;
    test(i);
    unlock(mutex);
    down(s[i]);
}

```

```

void put_forks(int i){
    lock(mutex);
    state[i] = THINKING;
    test(LEFT);
    test(RIGHT);
    unlock(mutex);
}

```

```

void test(int i){
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING){
        state[i] = EATING;
        up(s[i]);
    }
}

```

s[i] is 1, so down will not block.  
The value of s[i] decrements by 1.

|           | P1 | P2 | P3 | P4 | P5 |
|-----------|----|----|----|----|----|
| state     | T  | T  | E  | T  | T  |
| semaphore | 0  | 0  | 1  | 0  | 0  |



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    → down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    unlock(mutex);  
}
```

s[i] is 1, so down will not block.  
The value of s[i] decrements by 1.

```
void test(int i){  
    if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

|           | P1 | P2 | P3 | P4 | P5 |
|-----------|----|----|----|----|----|
| state     | T  | T  | E  | T  | T  |
| semaphore | 0  | 0  | 0  | 0  | 0  |



```

void philosopher(int i){
    while(TRUE){
        think();
        take_forks(i);
        eat();
        put_forks();
    }
}

```

```

void take_forks(int i){
    lock(mutex);
    state[i] = HUNGRY;
    test(i);
    unlock(mutex);
    down(s[i]);
}

```

```

void put_forks(int i){
    lock(mutex);
    state[i] = THINKING;
    test(LEFT);
    test(RIGHT)
    unlock(mutex);
}

```

```

void test(int i){
    if (state[i] = HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING){
        state[i] = EATING;
        up(s[i]);
    }
}

```

|           | P1 | P2 | P3 | P4 | P5 |
|-----------|----|----|----|----|----|
| state     | T  | T  | E  | H  | T  |
| semaphore | 0  | 0  | 0  | 0  | 0  |

blocked

# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT)  
    unlock(mutex);  
}  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

|           | P1 | P2 | P3 | P4 | P5 |
|-----------|----|----|----|----|----|
| state     | T  | T  | T  | H  | T  |
| semaphore | 0  | 0  | 0  | 0  | 0  |



# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}  
→
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT)  
    unlock(mutex);  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
} →
```

|           | P1 | P2 | P3 | P4 | P5 |
|-----------|----|----|----|----|----|
| state     | T  | T  | T  | E  | T  |
| semaphore | 0  | 0  | 0  | 0  | 0  |



**UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    unlock(mutex);  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

|           | P1 | P2 | P3 | P4 | P5 |
|-----------|----|----|----|----|----|
| state     | T  | T  | T  | E  | T  |
| semaphore | 0  | 0  | 0  | 1  | 0  |

blocked

# Solution to Dining Philosophers

```
void philosopher(int i){  
    while(TRUE){  
        think();  
        take_forks(i);  
        → eat();  
        put_forks();  
    }  
}
```

```
void take_forks(int i){  
    lock(mutex);  
    state[i] = HUNGRY;  
    test(i);  
    unlock(mutex);  
    down(s[i]);  
}
```

```
void put_forks(int i){  
    lock(mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT)  
    unlock(mutex);  
}
```

```
void test(int i){  
    if (state[i] = HUNGRY &&  
        state[LEFT] != EATING &&  
        state[RIGHT] != EATING){  
        state[i] = EATING;  
        up(s[i]);  
    }  
}
```

| P1 | P2 | P3 | P4 | P5 |
|----|----|----|----|----|
| T  | T  | T  | E  | T  |
| 0  | 0  | 0  | 0  | 0  |

wakeup



# Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- *Abstract data type*, internal variables only accessible by code within the procedure
- Only one process may be active within the monitor at a time
- But not powerful enough to model some synchronization schemes

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }

    procedure Pn (...) {.....}

    Initialization code (...) { ... }
}
```

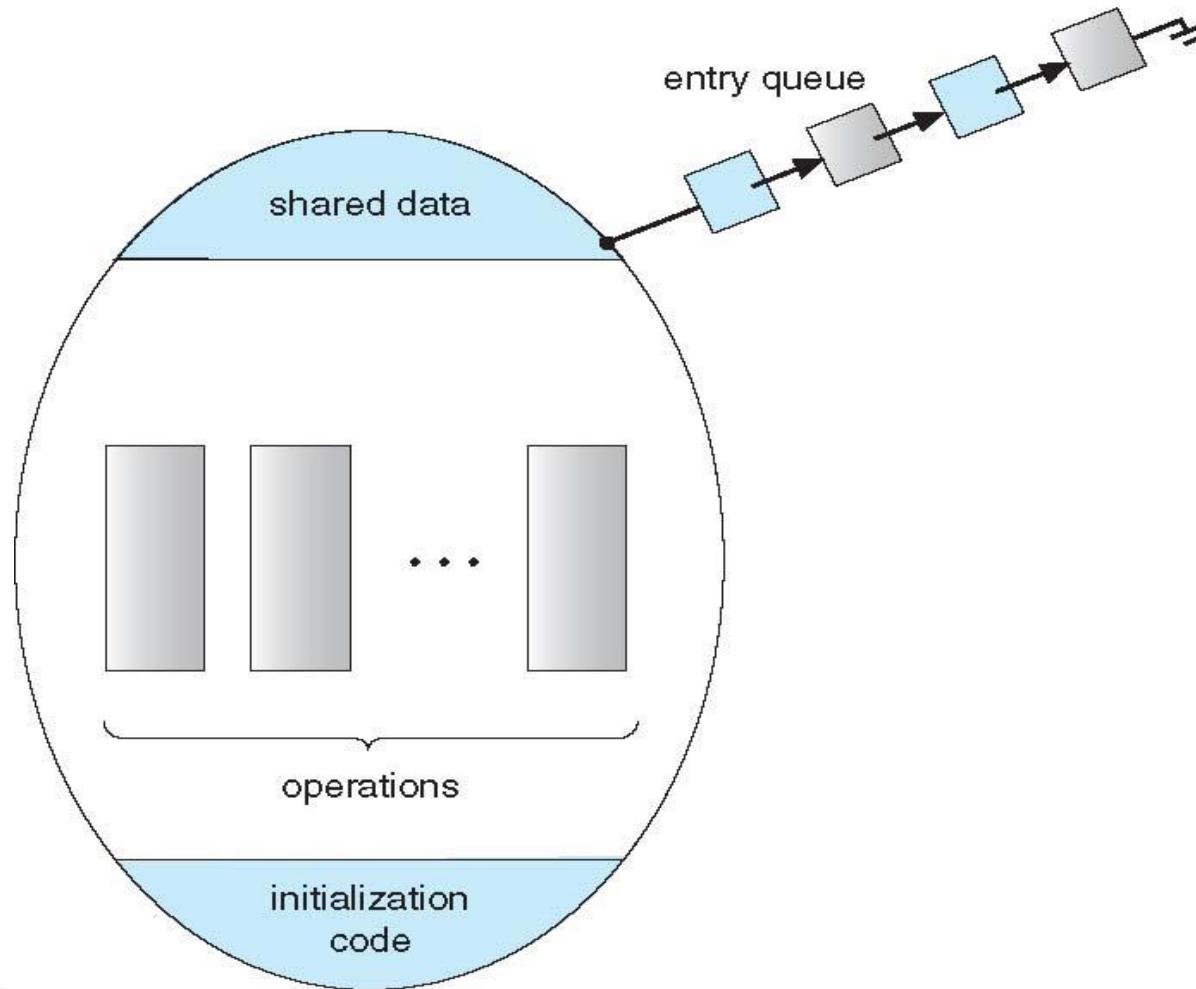


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Schematic view of a Monitor



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Condition Variables

- `condition x, y;`
- Two operations are allowed on a condition variable:
  - `x.wait()` – a process that invokes the operation is suspended until `x.signal()`
  - `x.signal()` – resumes one of processes (if any) that invoked `x.wait()`
    - If no `x.wait()` on the variable, then it has no effect on the variable

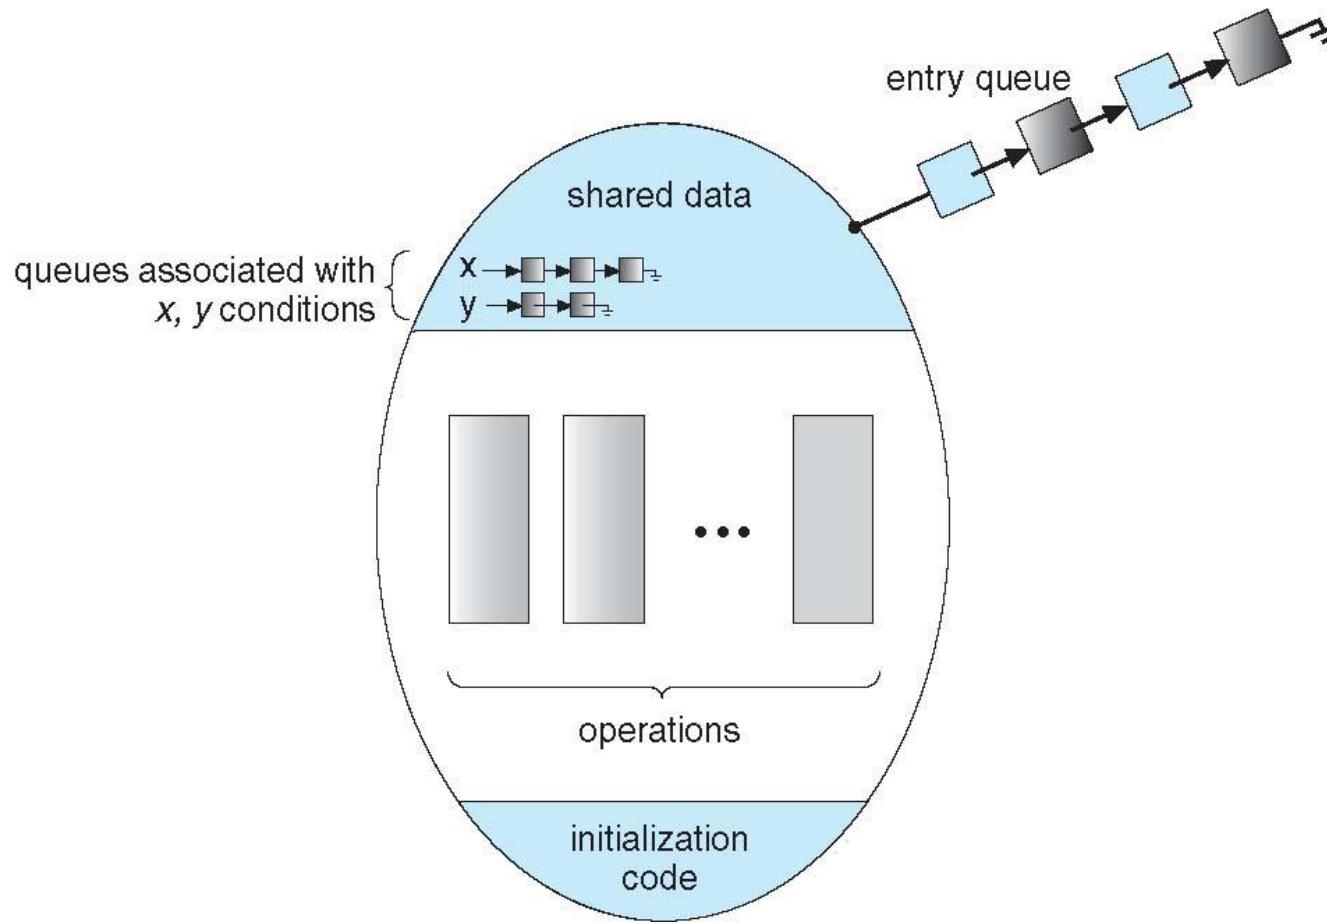


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Monitor with Condition Variables



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Condition Variables Choices

- If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?
  - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
  - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
  - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
  - Both have pros and cons – language implementer can decide
  - Monitors implemented in Concurrent Pascal compromise
    - P executing signal immediately leaves the monitor, Q is resumed
    - Implemented in other languages including Mesa, C#, Java



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING; HUNGRY, EATING) state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```



# Solution to Dining Philosophers (Cont.)

```
void test (int i) {  
    if ((state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}  
  
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}
```

# Solution to Dining Philosophers (Cont.)

- Each philosopher  $i$  invokes the operations `pickup()` and `putdown()` in the following sequence:

`DiningPhilosophers.pickup(i) ;`

**EAT**

`DiningPhilosophers.putdown(i) ;`

- No deadlock, but starvation is possible



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Deadlocks



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Contents

- The Deadlock Problem
- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
  - Deadlock Prevention
  - Deadlock Avoidance – Banker's algorithm
  - Deadlock Detection
  - Recovery from Deadlock



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.
- Example 1
  - System has 2 disk drives.
  - $P_1$  and  $P_2$  each hold one disk drive and each needs another one.
- Example 2
  - semaphores  $A$  and  $B$ , initialized to 1

$P_0$

wait (A);  
wait (B);

$P_1$

wait(B);  
wait(A);

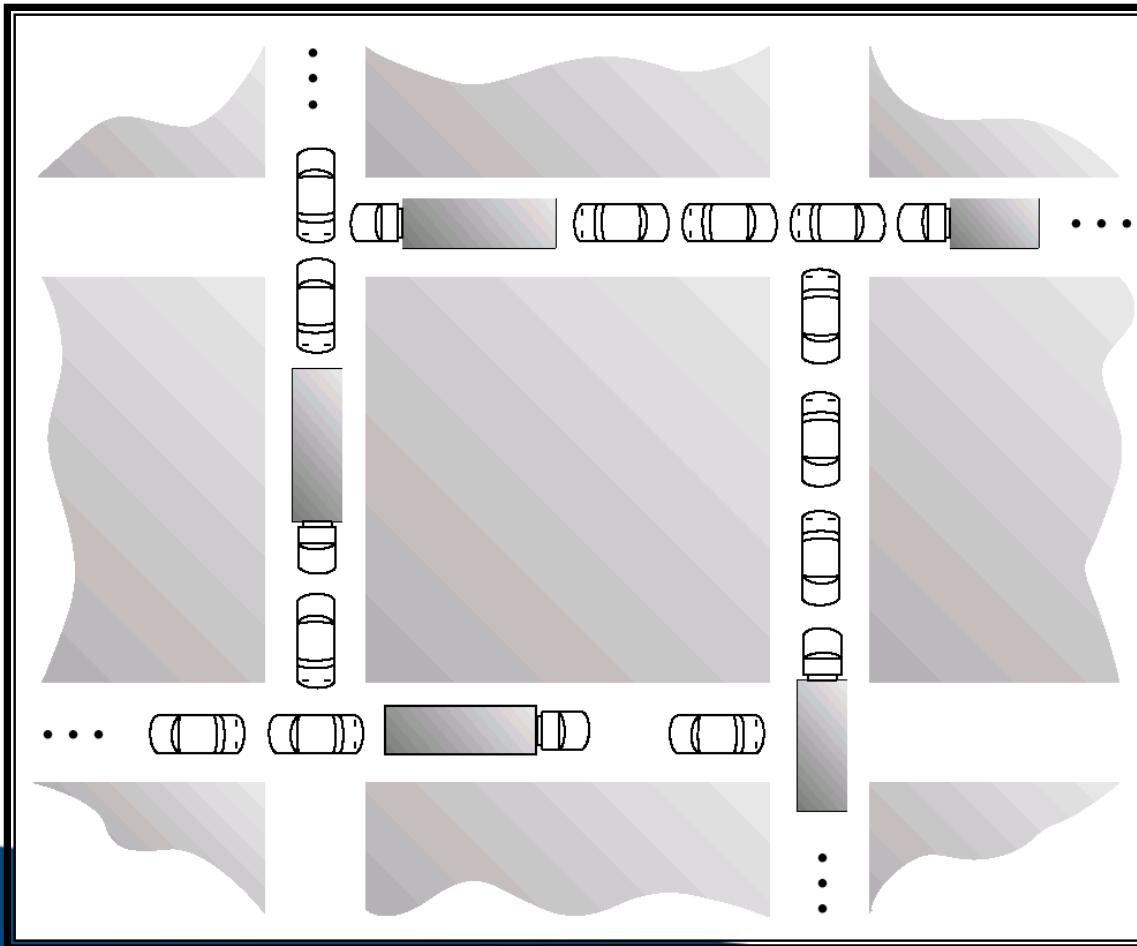


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Traffic Deadlock Example

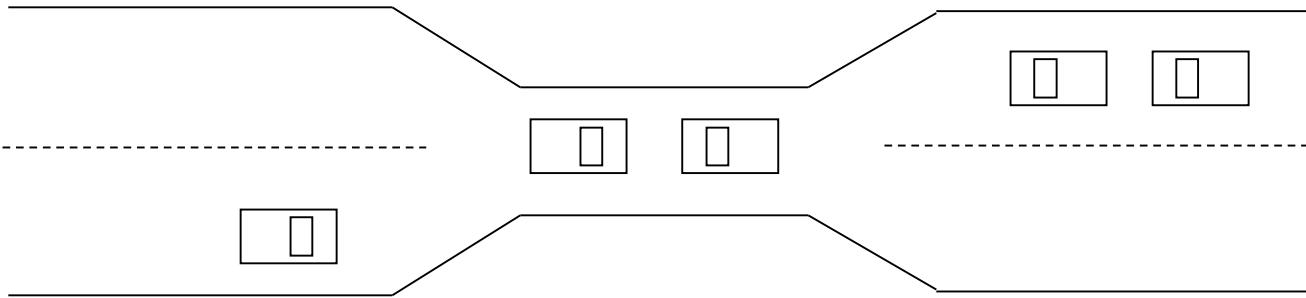


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Bridge Crossing Example



- Traffic only in one direction.
- Each section of a bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation is possible.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# System Model

- Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource as follows:
  - request
  - use
  - release

# Deadlock Characterization

Deadlock can arise if four conditions hold **simultaneously**.

- **Mutual exclusion**: only one process at a time can use a resource.
- **Hold and wait**: a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption**: a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait**: there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .

- $V$  is partitioned into two types:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the **processes** in the system.
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all **resource** types in the system.
- $E$  is partitioned into two types:
  - **request** edge – directed edge  $P_1 \rightarrow R_j$
  - **assignment** edge – directed edge  $R_j \rightarrow P_i$



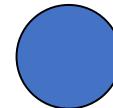
**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013

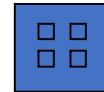


# Resource-Allocation Graph (Cont.)

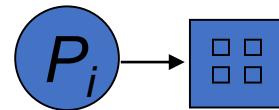
- Process



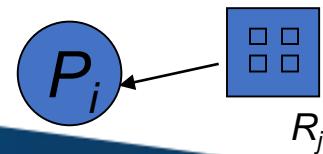
- Resource type with 4 instances



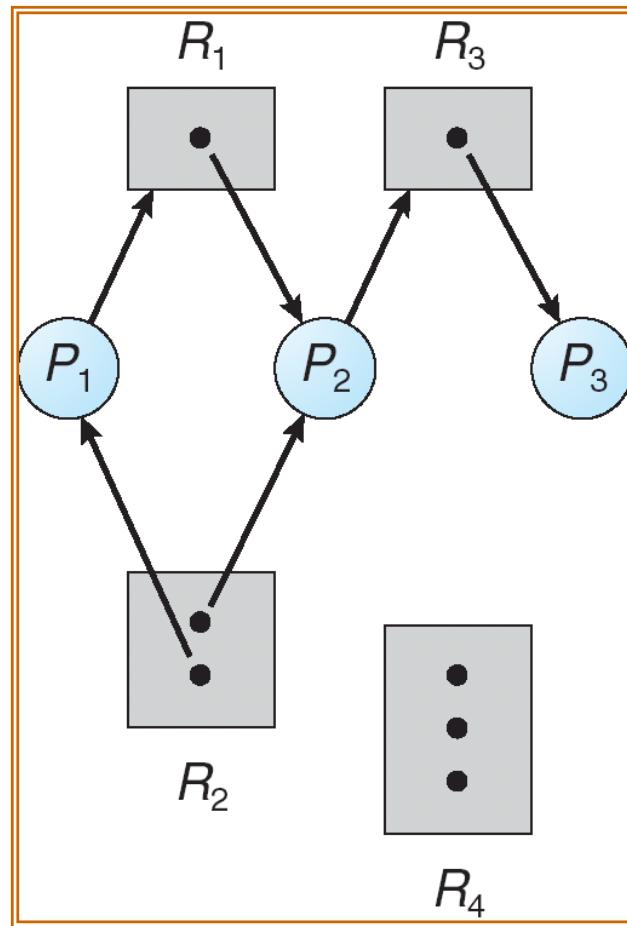
- $P_i$  requests an instance of  $R_j$



- $P_i$  is holding an instance of  $R_j$



# Example of a Resource Allocation Graph

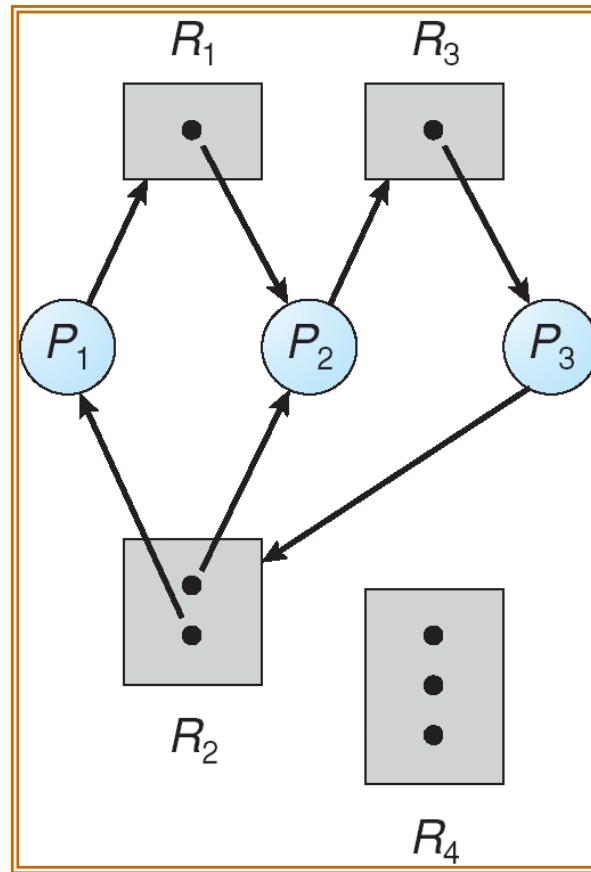


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Resource Allocation Graph with a Deadlock

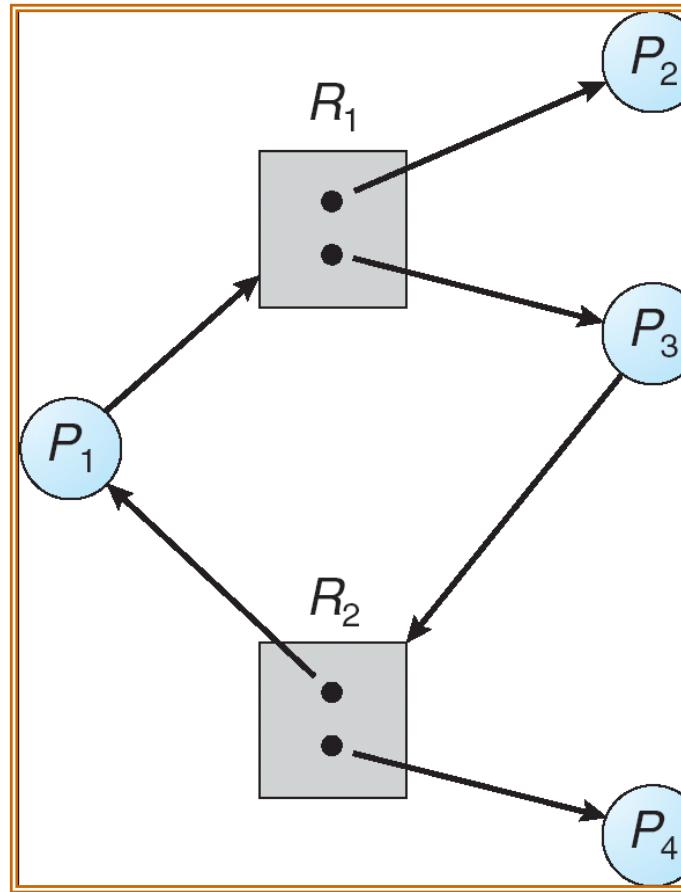


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Graph with a Cycle But No Deadlock



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



PRESIDENCY GROUP  
OVER  
**40**  
YEARS  
OF ACADEMIC  
WISDOM

# Basic Facts

- If graph contains no cycles  $\Rightarrow$  no deadlock.
- If graph contains a cycle
  - if only one instance per resource type, then deadlock.
  - if several instances per resource type, possibility of deadlock.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state.
  - deadlock *prevention*
  - deadlock *avoidance*
- Allow the system to enter a deadlock state and then recover.
  - deadlock *detection* and *recovery*
- Ignore the problem and pretend that deadlocks never occur in the system
  - used by most operating systems, including UNIX and Linux.

# Deadlock Prevention

Restrain the ways request can be made.

- Mutual Exclusion
  - not required for sharable resources  
e.g. read-only files
  - must hold for nonsharable resources  
e.g. printers

# Deadlock Prevention

- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources.
  - require process to request and be allocated all its resources before it begins execution, or
  - allow process to request resources only when the process has none.
- side effects
  - low resource utilization
  - starvation possible



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Deadlock Prevention

- No Preemption
  - Implicit preemption
    - If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
    - Preempted resources are added to the list of resources for which the process is waiting.
    - Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting
  - suitable for resources whose state can be easily saved and restored, such as CPU registers and memory space, but generally cannot apply to resources like printers and tape drives



# Deadlock Prevention

- Circular Wait
  - impose a total ordering of all resource types
  - require that each process requests resources in an increasing order of enumeration, or
    - whenever a process requests a resource, it has released any resources of higher order



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Deadlock Avoidance

Requires that the system has some additional *a priori* information available.

- Simplest and most useful model requires that each process declare the **maximum number of resources of each type** that it may need.
- The **deadlock-avoidance algorithm** dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
- **Resource-allocation state** is defined by the number of available and allocated resources, and the maximum demands of the processes.



# Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a **safe state**.
- The system is in safe state if there exists a **safe sequence** of *all* processes in the systems.
- A sequence  $\langle P_1, P_2, \dots, P_n \rangle$  is safe if for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$ .
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished.
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate.
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Example

- A system with 12 tape drives
- Current state

|                | Max Need | Current Hold |
|----------------|----------|--------------|
| P <sub>0</sub> | 10       | 5            |
| P <sub>1</sub> | 4        | 2            |
| P <sub>2</sub> | 9        | 2            |

Safe sequence: <P1, P0, P2>

It is a safe state.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks.
- If a system is in unsafe state  $\Rightarrow$  **possibility** of deadlock.
- **Deadlock avoidance**  $\Rightarrow$  ensure that a system will never enter an unsafe state.

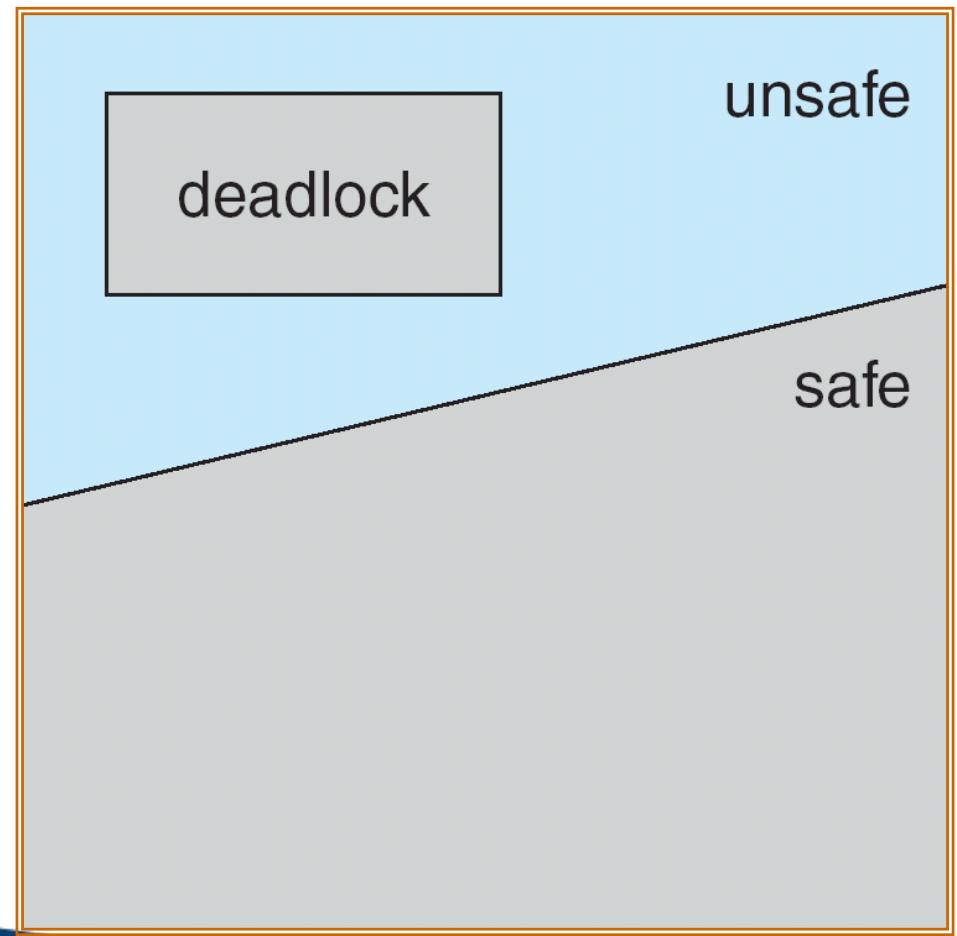


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Safe, Unsafe , Deadlock State



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



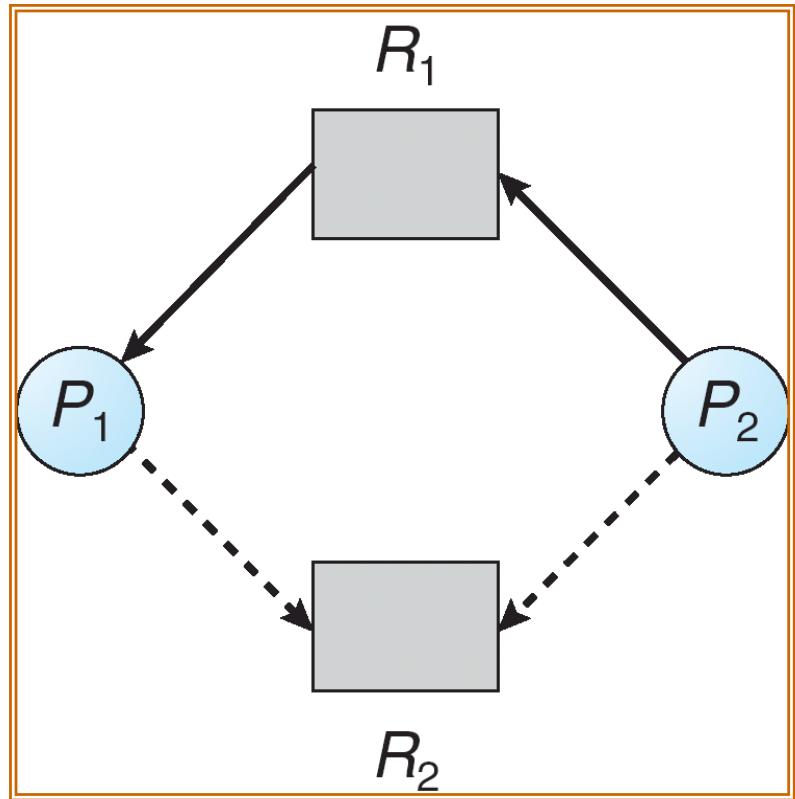
# Avoidance algorithms

- Single instance of a resource type
  - Use a **resource-allocation graph**
- Multiple instances of a resource type
  - Use the **banker's algorithm**

# Resource-Allocation Graph Scheme

- **Claim edge**  $P_i \rightarrow R_j$  indicated that process  $P_j$  may request resource  $R_j$ ; represented by a dashed line.
- Claim edge converts to **request edge** when a process requests a resource.
- Request edge converted to an **assignment edge** when the resource is allocated to the process.
- When a resource is released by a process, assignment edge reconverts to a claim edge.
- Resources must be claimed *a priori* in the system.

# Resource-Allocation Graph



if  $P_2$  request  $R_2$ ,  
and  $R_2$  is currently free,  
but ...

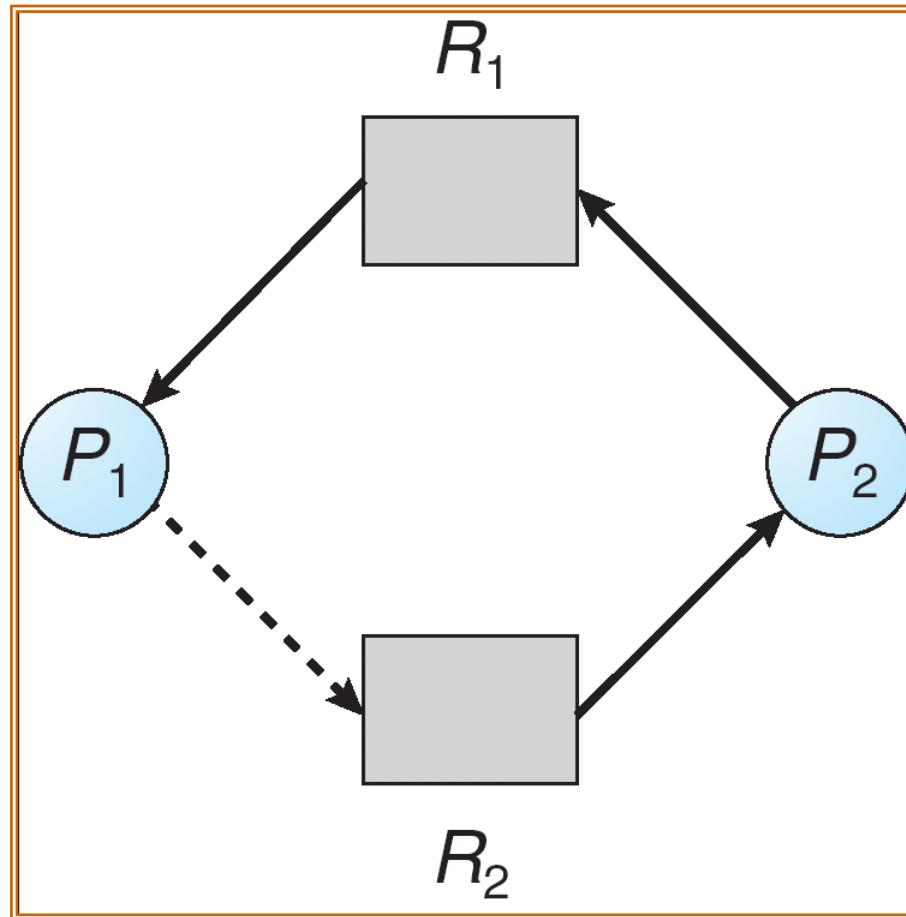


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Unsafe State In Resource-Allocation Graph



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

# Banker's Algorithm

- Suitable for multiple instances.
- Each process must **a priori claim** maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.
- Basic idea: keep the system in **safe state**



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- **Available:** Vector of length  $m$ . If  $\text{available}[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
- **Max:**  $n \times m$  matrix. If  $\text{Max}[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
- **Allocation:**  $n \times m$  matrix. If  $\text{Allocation}[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .
- **Need:**  $n \times m$  matrix. If  $\text{Need}[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.

$$\text{Need } [i,j] = \text{Max}[i,j] - \text{Allocation } [i,j].$$

# Safety Algorithm

1. Let **Work** and **Finish** be vectors of length m and n, respectively. Initialize:

Work = Available

Finish [i] = false for  $i = 0, 1, \dots, n - 1$ .

2. Find an  $i$  such that both:

(a) Finish [i] = false

(b)  $\text{Need}_i \leq \text{Work}$

If no such  $i$  exists, go to step 4.

3.  $\text{Work} = \text{Work} + \text{Allocation}_i$

Finish[i] = true

go to step 2.

4. If Finish [i] == true for all  $i$ , then the system is in a safe state.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Resource-Request Algorithm for Process $P_i$

$Request$  = request vector for process  $P_i$ . If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$ .

1. If  $Request_i \leq Need_i$  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If  $Request_i \leq Available$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available.
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- | If **safe**  $\Rightarrow$  the resources are **allocated** to  $P_i$ .
- | If **unsafe**  $\Rightarrow P_i$  must **wait**, and the old resource-allocation state is restored



# *Example of Banker's Algorithm*

- 5 processes  $P_0$  through  $P_4$
- 3 resource types:  
 $A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances).
- Snapshot at time  $T_0$ :

|       | <u>Allocation</u> | <u>Max</u>  | <u>Available</u> | <u>Need</u> |
|-------|-------------------|-------------|------------------|-------------|
|       | $A \ B \ C$       | $A \ B \ C$ | $A \ B \ C$      | $A \ B \ C$ |
| $P_0$ | 0 1 0             | 7 5 3       | 3 3 2            | 7 4 3       |
| $P_1$ | 2 0 0             | 3 2 2       |                  | 1 2 2       |
| $P_2$ | 3 0 2             | 9 0 2       |                  | 6 0 0       |
| $P_3$ | 2 1 1             | 2 2 2       |                  | 0 1 1       |
| $P_4$ | 0 0 2             | 4 3 3       |                  | 4 3 1       |



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Example (Cont.)

- The content of the matrix *Need* is defined to be *Max – Allocation*.

| <u>Need</u> |       |
|-------------|-------|
|             | A B C |
| $P_0$       | 7 4 3 |
| $P_1$       | 1 2 2 |
| $P_2$       | 6 0 0 |
| $P_3$       | 0 1 1 |
| $P_4$       | 4 3 1 |

The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria.

$m=3, n=5$

Step 1 of Safety Algo

Work = Available

Work = 

|   |   |   |   |   |
|---|---|---|---|---|
| 3 | 3 | 2 |   |   |
| 0 | 1 | 2 | 3 | 4 |

Finish = 

|       |       |       |       |       |
|-------|-------|-------|-------|-------|
| false | false | false | false | false |
|-------|-------|-------|-------|-------|

For  $i = 0$

Need<sub>0</sub> = 7, 4, 3

Finish [0] is false and Need<sub>0</sub> > Work

So P<sub>0</sub> must wait

Step 2:

But Need  $\leq$  Work

For  $i = 1$

Need<sub>1</sub> = 1, 2, 2

Step 2:

1,2,2      3,3,2

Finish [1] is false and Need<sub>1</sub> < Work

So P<sub>1</sub> must be kept in safe sequence

3, 3, 2      2, 0, 0

Step 3

Work = Work + Allocation<sub>1</sub>

Work = 

|   |   |   |   |   |
|---|---|---|---|---|
| 5 | 3 | 2 |   |   |
| 0 | 1 | 2 | 3 | 4 |

Finish = 

|       |      |       |       |       |
|-------|------|-------|-------|-------|
| false | true | false | false | false |
|-------|------|-------|-------|-------|

For  $i = 2$

Need<sub>2</sub> = 6, 0, 0

Step 2:

6, 0, 0      5, 3, 2

Finish [2] is false and Need<sub>2</sub> > Work

So P<sub>2</sub> must wait

For  $i=3$

Need<sub>3</sub> = 0, 1, 1

Step 2:

Finish [3] = false and Need<sub>3</sub> < Work

So P<sub>3</sub> must be kept in safe sequence

5, 3, 2      2, 1, 1

Step 3

Work = Work + Allocation<sub>3</sub>

Work = 

|   |   |   |   |   |
|---|---|---|---|---|
| 7 | 4 | 3 |   |   |
| 0 | 1 | 2 | 3 | 4 |

Finish = 

|       |      |       |      |       |
|-------|------|-------|------|-------|
| false | true | false | true | false |
|-------|------|-------|------|-------|

For  $i = 4$

Need<sub>4</sub> = 4, 3, 1

Step 2:

4, 3, 1      7, 4, 3

Finish [4] = false and Need<sub>4</sub> < Work

So P<sub>4</sub> must be kept in safe sequence

7, 4, 3      0, 0, 2

Step 3

Work = Work + Allocation<sub>4</sub>

Work = 

|   |   |   |   |   |
|---|---|---|---|---|
| 7 | 4 | 5 |   |   |
| 0 | 1 | 2 | 3 | 4 |

Finish = 

|       |      |       |      |      |
|-------|------|-------|------|------|
| false | true | false | true | true |
|-------|------|-------|------|------|

For  $i = 0$

Need<sub>0</sub> = 7, 4, 3

Step 2:

7, 4, 3      7, 4, 5

Finish [0] is false and Need  $<$  Work

So P<sub>0</sub> must be kept in safe sequence

7, 4, 5

0, 1, 0

Step 3

Work = Work + Allocation<sub>0</sub>

A    B    C  
7    5    5

0    1    2    3    4

Finish = 

|      |      |       |      |      |
|------|------|-------|------|------|
| true | true | false | true | true |
|------|------|-------|------|------|

For  $i = 2$

Need<sub>2</sub> = 6, 0, 0

Step 2:

6, 0, 0      7, 5, 5

Finish [2] is false and Need<sub>2</sub> < Work

So P<sub>2</sub> must be kept in safe sequence

7, 5, 5

3, 0, 2

Step 3

Work = Work + Allocation<sub>2</sub>

A    B    C  
10    5    7

0    1    2    3    4

Finish = 

|      |      |      |      |      |
|------|------|------|------|------|
| true | true | true | true | true |
|------|------|------|------|------|

Finish [i] = true for  $0 \leq i \leq n$

Step 4

Hence the system is in Safe state

The safe sequence is P<sub>1</sub>, P<sub>3</sub>, P<sub>4</sub>, P<sub>0</sub>, P<sub>2</sub>



PRESIDENCY  
UNIVERSITY

Private University Estd. in Karnataka State by Act No. 41 of 2013

PRESIDENCY GROUP  
OVER  
40  
YEARS OF ACADEMIC  
WISDOM

# Example: $P_1$ Request (1,0,2)

- Check that Request  $\leq$  Available
  - that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true.
- Pretend the request has been granted

|       | <u>Allocation</u> | <u>Need</u> | <u>Available</u> |
|-------|-------------------|-------------|------------------|
|       | A B C             | A B C       | A B C            |
| $P_0$ | 0 1 0             | 7 4 3       | 2 3 0            |
| $P_1$ | 3 0 2             | 0 2 0       |                  |
| $P_2$ | 3 0 2             | 6 0 0       |                  |
| $P_3$ | 2 1 1             | 0 1 1       |                  |
| $P_4$ | 0 0 2             | 4 3 1       |                  |

- Executing safety algorithm shows that sequence  $< P_1, P_3, P_4, P_0, P_2 >$  satisfies safety requirement.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Example (cont.)

- Can request for (3,3,0) by  $P_4$  be granted?
  - No, because the resources are not available.
- Can request for (0,2,0) by  $P_0$  be granted?
  - No. Even though the resources are available, the resulting state is **unsafe**

|       | <u>Allocation</u> |          |          | <u>Need</u> |          |          | <u>Available</u> |          |          |
|-------|-------------------|----------|----------|-------------|----------|----------|------------------|----------|----------|
|       | <i>A</i>          | <i>B</i> | <i>C</i> | <i>A</i>    | <i>B</i> | <i>C</i> | <i>A</i>         | <i>B</i> | <i>C</i> |
| $P_0$ | 0                 | 3        | 0        | 7           | 2        | 3        | 2                | 1        | 0        |
| $P_1$ | 3                 | 0        | 2        | 0           | 2        | 0        |                  |          |          |
| $P_2$ | 3                 | 0        | 2        | 6           | 0        | 0        |                  |          |          |
| $P_3$ | 2                 | 1        | 1        | 0           | 1        | 1        |                  |          |          |
| $P_4$ | 0                 | 0        | 2        | 4           | 3        | 1        |                  |          |          |



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Single Instance of Each Resource Type

- Maintain *wait-for* graph
  - Nodes are processes.
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock.
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph.

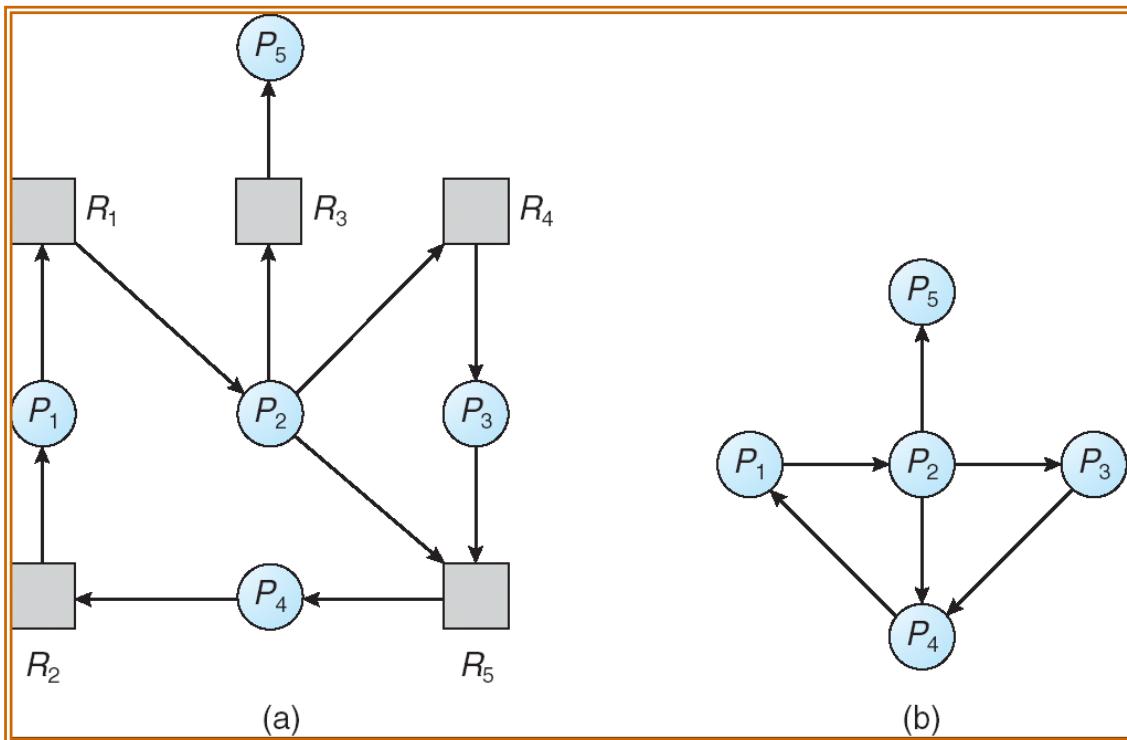


**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding **wait-for** graph

# Several Instances of a Resource Type

- **Available:** A vector of length  $m$  indicates the number of available resources of each type.
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $\text{Request}[i_j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Detection Algorithm

1. Let  $Work$  and  $Finish$  be vectors of length  $m$  and  $n$ , respectively

Initialize:

- (a)  $Work = Available$
- (b) For  $i = 1, 2, \dots, n$ , if  $Allocation_i \neq 0$ , then  
 $Finish[i] = \text{false}$ ; otherwise,  $Finish[i] = \text{true}$ .

2. Find an index  $i$  such that both:

- (a)  $Finish[i] == \text{false}$
- (b)  $Request_i \leq Work$

If no such  $i$  exists, go to step 4.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Detection Algorithm (Cont.)

3.  $Work = Work + Allocation_i$

$Finish[i] = true$

go to step 2.

4. If  $Finish[i] == \text{false}$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  $Finish[i] == \text{false}$ , then  $P_i$  is deadlocked.

The algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state.



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ;
- Three resource types:  
A (7 instances), B (2 instances), and C (6 instances).
- Snapshot at time  $T_0$ :

|       | <u>Allocation</u> |   |   | <u>Request</u> |   |   | <u>Available</u> |   |   |
|-------|-------------------|---|---|----------------|---|---|------------------|---|---|
|       | A                 | B | C | A              | B | C | A                | B | C |
| $P_0$ | 0                 | 1 | 0 | 0              | 0 | 0 | 0                | 0 | 0 |
| $P_1$ | 2                 | 0 | 0 | 2              | 0 | 2 |                  |   |   |
| $P_2$ | 3                 | 0 | 3 | 0              | 0 | 0 |                  |   |   |
| $P_3$ | 2                 | 1 | 1 | 1              | 0 | 0 |                  |   |   |
| $P_4$ | 0                 | 0 | 2 | 0              | 0 | 2 |                  |   |   |

- Sequence  $\langle P_0, P_2, P_3, P_4, P_1 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$ .



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Example (Cont.)

- Suppose  $P_2$  requests an additional instance of type C , make the *Request* matrix to be

|       | <u>Request</u> |   |   |
|-------|----------------|---|---|
|       | A              | B | C |
| $P_0$ | 0              | 0 | 0 |
| $P_1$ | 2              | 0 | 1 |
| $P_2$ | 0              | 0 | 1 |
| $P_3$ | 1              | 0 | 0 |
| $P_4$ | 0              | 0 | 2 |

- State of system?
  - Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes' requests.
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$ .



# Detection-Algorithm Usage

- When, and how often, to invoke depends on:
  - **How often** a deadlock is likely to occur?
  - **How many** processes will need to be rolled back?
    - one for each disjoint cycle
- called at specific interval, say every an hour
- called whenever a request cannot be granted immediately
  - can locate the specific process “caused” the deadlock
- called when CPU utilization drops, say, below 40%
  - If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph, and we cannot tell which process “caused” the deadlock.



# Recovery from Deadlock

## Process Termination

- Abort all deadlocked processes.
- Abort one process at a time until the deadlock cycle is eliminated.
- In which order should we choose to abort?
  - Priority of the process.
  - How long process has computed, and how much longer to completion.
  - Resources the process has used.
  - Resources process needs to complete.
  - How many processes will need to be terminated.
  - Is process interactive or batch?



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Resource Preemption

- **Selecting a victim** – minimize cost
  - factors similar as above mentioned
- **Rollback** – return to some safe state, restart process for that state
  - OS should keep track of states of all processes
- **Starvation** – same process may always be picked as victim
  - solution: include number of rollback in cost factor



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013



# Thank You



**PRESIDENCY  
UNIVERSITY**

Private University Estd. in Karnataka State by Act No. 41 of 2013

